


SQL Injection Attacks:

Detection in a Web Application Environment



NETWORKS

Table of Contents

1	Foreword.....	1
2	Background	2
2.1	Web Application Environment.....	2
2.2	SQL Attack Overview	2
2.3	Applications Open to SQL Injection	3
3	The challenge with detection.....	4
3.1	Effective Security	4
3.2	Types of attacks	5
3.2.1	Tautologies	5
3.2.2	Union Query.....	5
3.2.3	Illegal/Logically Incorrect Queries	6
3.2.4	Stored Procedure Attacks.....	6
3.2.5	Alternate Encoding Obfuscation	7
3.2.6	Combination Attacks	7
3.3	Detection at the Web Tier	7
3.3.1	Detecting SQL Injection Challenges	7
3.3.2	Web Tier Detection in Practice	9
3.4	A better way – a Database Firewall.....	9
4	Conclusions	11

1 Foreword

It has been nearly two decades since the original research paper on SQL Injection was published. Over the years the SQL Injection threat has grown to the point where now we are seeing far more devastating SQL Injection attacks than ever before. Organizations are being breached via SQL Injection attacks that slip seamlessly through the network firewall over port 80 (HTTP) or 443 (SSL) and bypass their web application firewalls (WAF) through obfuscation. At that point the attacker can exploit the soft internal networks and vulnerable databases. It's a well-know threat and occupies the number one spot on OWASP top ten threats year after year.

Detecting SQL fragments injected into a Web application has proven extremely challenging. There are several tacks enterprises can take – prevention, remediation, and mitigation. When implementing prevention and remediation efforts, the enterprise strives to develop secure code and/or encrypt confidential data stored in the database. However, these are not always available options. For Example, in some cases the application source code may have been developed by a third party and not be available for modification. Additionally, patching deployed code requires significant resources and time. Therefore rewriting an existing operational application would need to be prioritized ahead of projects driving new business. Similarly, efforts to encrypt confidential data stored in the database can take even longer and require more resources. Given today's compressed development cycles, and limited number of developers with security domain experience, even getting the code rewrite project off the ground could prove difficult.

One approach often employed in an attempt to identify SQL injection attacks is a WAF. A WAF operates in front of the Web server and monitors the traffic into and out of the Web servers. A WAF attempts to identify patterns that constitute a threat (see Figure 1). While this can be effective in detecting certain classes of attacks against Web applications, it has proven ineffective in detecting all but the simplest SQL injection attacks.

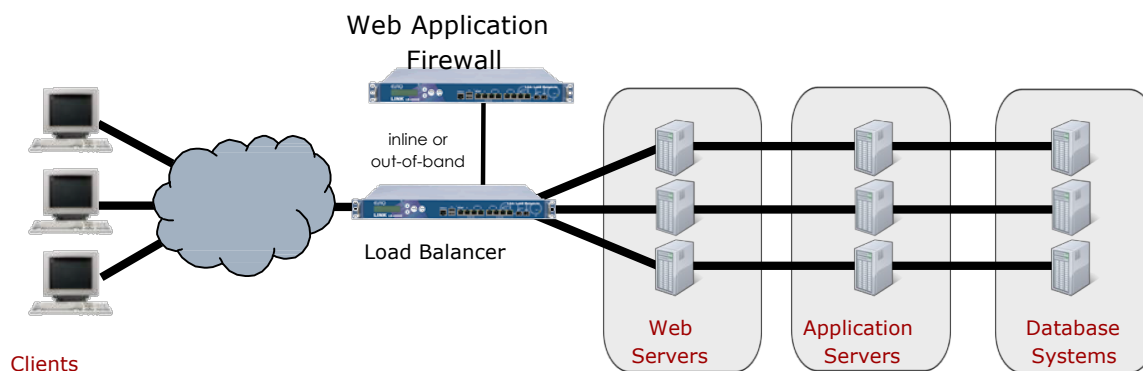


Figure 1 - Network Placement of a Web Application Firewall

This shouldn't be taken to infer a WAF isn't a useful component within a Web security environment. To the contrary, WAFs provide a number of benefits including reasonable protection from header injection and XSS attacks. A WAF should always be considered as part of a Web security defense in depth strategy. However, for SQL injection prevention organizations are turning to database security appliances. To better understand why WAFs are a relatively ineffective for SQL Injection simply conduct a Web search for "WAF bypass".

2 Background

2.1 Web Application Environment

Before we drive into a discussion on the approaches to detect and prevent SQL injection attacks, let's first explore the Web application environment. Web application information is presented to the Web server by the user's client, in the form of URL's, cookies and form inputs (POSTs and GETs). These inputs drive both the logic of the application and the queries those applications create and send to a database to extract relevant data.

Unfortunately, many applications do not adequately validate user input and are thus susceptible to SQL injection. Attackers capitalize on these flaws to attempt to cause the backend database to do something different than what the application (and the organization) intended. This can include extracting sensitive information, destroying information or executing a denial of service (DoS) attack that limits others' use of the application.

2.2 SQL Injection Attack Overview

SQL injection attacks are initiated by manipulating the data input on a Web form such that fragments of SQL instructions are passed to the Web application. The Web application then combines these rogue SQL fragments with the proper SQL dynamically generated by the application, thus creating valid SQL requests. These new, unanticipated requests cause the database to perform the task the attacker intends.

To clarify, consider the following simple example. Assume we have an application whose Web page contains a simple form with input fields for username and password. With these credentials the user can get a list of all credit card accounts they hold with a bank. Further assume that the bank's application was built with no consideration of SQL injection attacks.

As such, it is reasonable to assume the application merely takes the input the user types and places it directly into an SQL query constructed to retrieve that user's information. In PHP that query string would look something like this:

```
$query = "select accountName, accountNumber from  
creditCardAccounts where username='". $_POST["username"]. "'  
and password='". $_POST["password"]. "'"
```

Normally this would work properly as a user entered their credentials, say johnSmith and myPassword, and formed the query:

```
$query = "select accountName, accountNumber from  
creditCardAccounts where username='johnSmith' and  
password='myPassword'
```

This query would return one or more accounts linked to Mr. Smith.

Now consider someone with a devious intent. This person decides they want to see if they can access the account information of one or more of the bank's customers. To accomplish this they enter the following credential into the form:

```
' or 1=1 -- and anyThingsAtAll
```

When this SQL fragment is inserted into the SQL query by the application it becomes:

```
$query = "select accountName, accountNumber from  
creditCardAccounts where username='' or 1=1 -- and  
password= anyThingsAtAll
```

The injection of the term, ' or 1=1 --, accomplishes two things. First, it causes the first term in the SQL statement to be true for all rows of the query; second, the -- causes the rest of the statement to be treated as a comment and, therefore, ignored during run time. The result is that all the credit cards in the database, up to the limit the Web page will list, are returned and the attacker has stolen the valuable information they were seeking.

It should be noted that this simple example is just one of literally an infinite number of variations that could be used to accomplish the same attack. Further, there are many other ways to exploit a vulnerable application. We will discuss more of these attacks as we delve into the efficacy of various attack mitigation techniques.

2.3 Applications Vulnerable to SQL Injection

There are a number of factors that conspire to make securely written applications a rarity. First, many applications were written at a time when Web security was not a major consideration. This is especially true of SQL injection. While recently SQL injection is being discussed at security conferences and other settings, the attack frequency of SQL injection only

five or so years ago was low enough that most developers were simply not aware.

In addition, the application may have been initially written as an internal application with a lower security threshold and subsequently exposed to the Web without considering the security ramifications. Even applications being written and deployed today often inadequately address security concerns. IBM's X-Force project recently found that 47% of all vulnerabilities that result in unauthorized disclosures are Web application vulnerabilities. Cross-Site Scripting & SQL injection vulnerabilities continue to dominate as the attack vector of choice.¹ Note that these reported vulnerabilities are for packaged applications from commercial software vendors. Vulnerabilities in custom applications were not reported. Since this software is generally not as carefully vetted for security robustness, it is reasonable to assume the problem is actually much bigger. According to Neira Jones, head of payment security for Barclays, 97% of data breaches worldwide are still due to an SQL injection somewhere along the line.²

Interestingly, modern environments and development approaches create a subtle vulnerability. With the advent of Web 2.0 there has been a shift in how developers treat user input. In these applications input is rarely provided by a simple form that directly transmits the information into the Web server for processing. In many cases, the JavaScript portion of the application performs input validation so the feedback to the user is handled more smoothly. This often creates the sense that the application is protected because of this very specific input validation; therefore, the validation on the server side is largely neglected. Unfortunately, attackers won't use the application to inject their input into the server component of the application. Rather, they leverage intermediate applications to capture the client-side input and allow them to manipulate it. Since the majority of the input validation is bypassed, the attacker can simply enter the SQL fragments needed to change the behavior of the database to accomplish their intent.

¹ IBM Internet Security Systems™ X-Force® Mid-Year Trend and Risk Report

² Techworld, "Barclays: 97% of data breaches still due to SQL injection"

3 The challenge with detection

3.1 Effective Security

The goal of any security technology is to provide a robust threat detection and avoidance mechanism that requires little or no setup, configuration or tuning. Further, if that technology relies on learning or training to determine what is normal or to improve its ability to detect threats, those learning periods must be short and well-defined. This is needed to expedite installation and minimize the risk of attacks contaminating the learned dataset. Keep in mind the longer the learning period, the more likely an attack will occur and the larger the dataset you need to review to ensure that an attack has not occurred. Finally, given that few Web applications remain static, effective protection must be easy to maintain in the face of on-going changes to the Web application.

3.2 Types of attacks

Previously a simple attack on a vulnerable application was described to illustrate how a SQL Injection attack can occur. The general class of attacks that the simple example falls into can be described as Tautological attacks. Tautologies are statements composed of simpler statements in such a way that makes the statement true regardless if simpler statements are true or false. For example, the statement "Either it will rain tomorrow or it will not rain tomorrow" is a tautology.

The complexity of detecting SQL injection can best be understood through a variety of examples demonstrating the various SQL injection attack classifications. This list is not exhaustive, but rather provides a sample of the most common injections seen in real deployments.

3.2.1 Tautologies

This attack works by inserting an "always true" statement into a WHERE clause of the SQL statement to extract data. These are often used in combination with the insertion of a -- to cause the remainder of a statement to be ignored ensuring extraction of largest amount of data. Tautological injections can include techniques to further mask SQL expression snippets, as demonstrated by the following example:

```
' or 'simple' like 'sim%' --  
' or 'simple' like 'sim' || 'ple' --
```

The || in the example is used to concatenate strings, when evaluated the text 'sim' || 'ple' becomes 'simple'.

3.2.2 Union Query

This attack exploits a vulnerable parameter by injecting a statement of the form:

```
foo'UNION SELECT <rest of injected query>
```

The attacker can insert any appropriate query to retrieve information from a table different from the one that was the target of the original statement. The database returns a dataset that is the union of the results of the original first query and the results of the injected second query.

3.2.3 Illegal/Logically Incorrect Queries

Attackers use this approach to gather important information about the type of database and its structure. Attacks of this nature are often used in the initial reconnaissance phase to gather critical knowledge used in other attacks. Returned error pages that are not filtered can be very instructive. Even if the application sanitizes error messages, the fact that an error is returned or not returned can reveal vulnerable or injectable parameters. Syntax errors identify injectable parameters; type errors help decipher data types of certain columns; logical errors, if returned to the user, can reveal table or column names.

The specific attacks within this class are largely the same as those used in a Tautological attack. The difference is that these are intended to determine how the system responds to different attacks by looking at the response to a normal input, an input with a logically true statement appended (typical tautological attack), an input with a logically false statement appended (to catch the response to failure) and an invalid statement to see how the system responds to bad SQL. This will often allow the attacker to see if an attack got through to the database even if the application does not allow the output from that statement to be displayed.

There are a myriad of examples. In fact, the attacker may initially use a bot to detect a vulnerable web site and then recursively use this class of attack forensically to learn application and database specifics.

The key point in listing this classification is that WAFs are unable to detect such attacks if the injections fall outside of the signatures created by the WAF learning process. As well, the WAF may not be exposed to error messages that the application (and a Database Firewall) will receive.

3.2.4 Stored Procedure Attacks

These attacks attempt to execute database stored procedures. The attacker initially determines the database type (potentially using illegal/logically incorrect queries) and then uses that knowledge to determine what stored procedures might exist. Contrary to popular belief using stored procedures does not make the database invulnerable to SQL injection attacks. Stored procedures can be susceptible to privilege escalation, buffer overflows, and even provide access to the operating system.

3.2.5 Alternate Encoding Obfuscation

In this case, text is injected as to avoid detection by defensive coding practices. It can also be very difficult to generate rules for a WAF to detect encoded input. Encodings, in fact, can be used in combination with other attack classifications. Since databases parse comments out of an SQL statement prior to processing it, comments are often used in the middle of an attack to hide the attack's pattern.

Scanning and detection techniques, including those used in WAFs, have not been effective against alternate encodings or comment based obfuscation because all possible encodings must be considered.

Note that these attacks may have no SQL keywords embedded as plain text, though it could run arbitrary SQL.

3.2.6 Combination Attacks

Many attack vectors may be employed in combination:

- learn information useful in generating additional successful injections (illegal/logically incorrect)
- gain access to systems other than the initial database accessed by the application (stored procedures)
- evade detection by masking intent of injection (alternate encoding)

3.3 Detection at the Web Tier

3.3.1 Detecting SQL Injection Challenges

Given the large variation in the form or pattern of SQL attacks, it can be very challenging to detect them from a point in front of the Web server. At this network location the Web Application Firewall is attempting to identify a possible snippet of SQL in the input stream of a Web application.

Why is it difficult to detect input injections at the Web tier? Remember, the WAF is not inspecting the SQL request as sent to the database by the application tier. Rather, it has URL's, cookies and form inputs (POSTs and GETs) to inspect. Inspecting each set of input values, a WAF must consider the wide range of acceptable input against what is considered unacceptable for each input field on each form.

Although many attacks use special characters that may not be expected in a typical form, two problems complicate detection. With no prior knowledge of the application it is not possible to know with certainty what characters are expected in any given field. Furthermore, in some cases the characters used do, in fact, occur in normal input and blocking them at the character level is not possible. Consider the single quote often used to delimit a string. Unfortunately, this character appears in names such as O'Brien or in possessive expressions like Steve's; therefore, single quotes are valid in some input fields.

As a result larger patterns must be considered, which are more demonstrative of an actual attack, to bring the false positives down to a reasonable rate. And this is where the problem begins. The choice then becomes: use a very general set of patterns such as checking for a single quote or the word "like" or possibly "or" to catch every conceivable attack or use a more complicated pattern that reduces the false positive rate.

Since there is a reasonable likelihood that general patterns exist in normal input, the WAF must then inspect all form input (in learning or training mode) for an extended period of time before it can determine which of these simple patterns can reliably be used to validate each form and each input field in the Web application. Considering the complexity, range and limited structure within the natural language used in forms, it can take a very long time to ensure that an adequate sample size has been gathered to confirm that selected detection patterns are not found in legitimate input. Complicating this further is the fact that some sections of an application are often used infrequently, extending even further the training time. An example would be business logic exercised according to the business cycle. Add it all up and you can see this approach requires an extensive time period to ensure that the learning cycle has adequately considered all the variations of valid input for each field on each form of the Web application.

Alternatively, as mentioned above, much more complex patterns that are clearly indicative of an attack can be used. Unfortunately, as we

demonstrated in our discussion of the attack types, the number and variation of possible attacks is so large that it is impossible to effectively cover all possible attack patterns. Creating the initial pattern set, keeping up with the evolving attacks and verifying that they are sufficiently unique as to not show up in some fields is an almost impossible task. And now, consider that the applications are also changing and evolving over time, requiring further, time-consuming learning.

3.3.2 Web Tier Detection in Practice

So how are WAF's used in the real world? One way is to use a combination of approaches, each aimed at reducing the negative effects of the other approach. These negative effects include limited capability to detect a SQL injection versus high number of false positives, complex configurations, and long training times. Specifically, a large set of patterns ranging from relatively simple to much more complex are used. Some patterns are configured to be applied to all input sources regardless of what is learned during training; some patterns are configured such that they will be removed, for a given input field, if they are contained within the training data. Some rules and patterns also attempt to classify the range of input by length and character set, for example, numerical fields.

The WAF is then placed into learning mode and allowed to learn until it is believed that a large enough set of each input field has been examined to reduce subsequent false positives. The resulting sets are then reviewed to determine if the learned set for some fields is considered too small, requiring additional learning time or manual manipulation. Other fields, whose default rule set have been reduced too far, are reviewed to determine what hand crafted rules can be constructed to increase the coverage.

This manual inspection process on top of the long learning cycle, while more effective than any one approach in isolation, is far from efficient. However, it still suffers the weaknesses of an administrator having to make decisions, configuring a significant number of rule/pattern sets for fields not effectively configured through training. This can be true even after a substantial learning period has been used.

This, in a nutshell, is why WAFs have been ineffective in curtailing SQL injection attacks. It's self evident, had WAFs been effective the size and scope of SQL injection attacks would not be increasing year over year.

3.4 A better way – a Database Firewall

Thus far we have described the method of detecting SQL injection attacks at the Web tier interface. A more effective and efficient method is to analyze the actual SQL generated by the application and presented to the database. The Database Firewall monitors the networks between the application servers and databases (see Figure 2). Why is this more effective and more efficient? The simple answer is that while the input into the Web tier has an enormous pattern set with very little structure associated with each input field, an application creates a comparatively small set of SQL statements (ignoring the literal values associated with those statements). In addition the structure of SQL statement lends themselves to structured analysis. Both of these factors make analysis more determinant than the rudimentary input pattern validation of a WAF. We will discuss how to deal with the variation of the literal values (the actual intended user input) below.

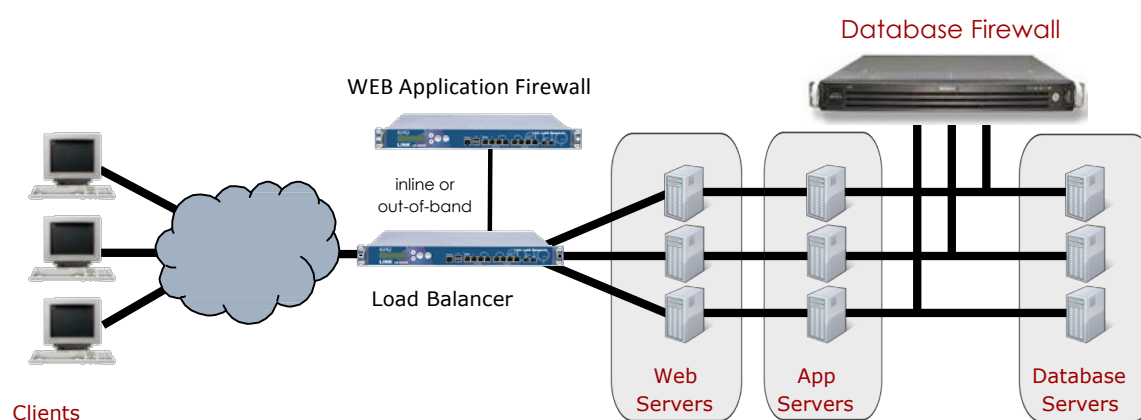


Figure 2 - Placement of Database Firewall

At the database interface, an SQL statement can be processed in much the same way the database itself processes it – breaking it down into the statement structure and separating out the literals. Once this is done the very first use of any given input will generate the unique SQL statements associated with that input – as opposed to needing a large sample set to determine what patterns are not present.

As a result the sample set for learning is already reduced from that required for a WAF to a much smaller set needed to train a device inspecting traffic between the application and database. Once a working training set is developed it can be used to analyze all subsequent SQL statements and any

structure differs from the known set can be immediately flagged. By inspecting traffic at the interface to the database, it is clear which commands are leveraging stored procedures and it is easy to analyze the strings passed to stored procedures to determine if they contain any attacks. Several techniques can be applied in this analysis, such as observing the lack of delimiting special characters within literal strings.

Although analyzing the stream of SQL statements as described above provides a significant improvement over a WAF sitting at the Web tier, a true Database Firewall requires additional capabilities.

As pointed out during the discussion about training a WAF, many of the input fields within an application may not be exercised often during normal operations. Fortunately, most modern applications build their SQL from a set of logic that operates much like a code generator. This fact means that, using a relatively small sample set, it is possible to construct a model of how an application builds statements. An Adaptive Database Firewall can then use that knowledge to analyze newly discovered statements and assess their likeliness of being an attack.

In addition, given the fact that an SQL injection attack must be constructed out of an existing statement in the application further simplifies the analysis. If a new statement can be created wholly by inserting a string into the literal field of an existing statement, then it becomes highly suspect. Combining these concepts provides a means of assessing any new statement using algorithms that determine:

- Uniqueness relative to other statements previously seen
- Ability for that statement to have been constructed from a previously known statement
- Likelihood that the statement could have been generated within the application itself

Although an Adaptive Database Firewall uses a number of other important algorithms for analyzing incoming SQL against the learned model (for each application), the three algorithms highlighted above demonstrate the substantial value of operating at the interface to the database. No other approach can come close to the accuracy provided with this architecture. Furthermore, no other solution can be deployed with as little configuration and as short a training interval.

4 Conclusions

The efficacy of a security solution is measured by the robustness of its threat detection and avoidance mechanisms, its ease of setup, configuration or tuning, and its ability to detect SQL injection attacks with low false positive rates. Using these measures a true Database Firewall is far superior to a WAF in detecting SQL injection attacks. This is true because an Adaptive Database Firewall can be trained quicker, has a lower inherent false positive levels and is capable of seeing through virtually all attack obfuscation techniques.

In the end, a multi-layer Web security strategy is the best solution, drawing on the strengths of all relevant technologies. Considering the seriousness of the SQL injection threat, an Adaptive Database Firewall should be a prominent element in every solution.